


Backup and recovery

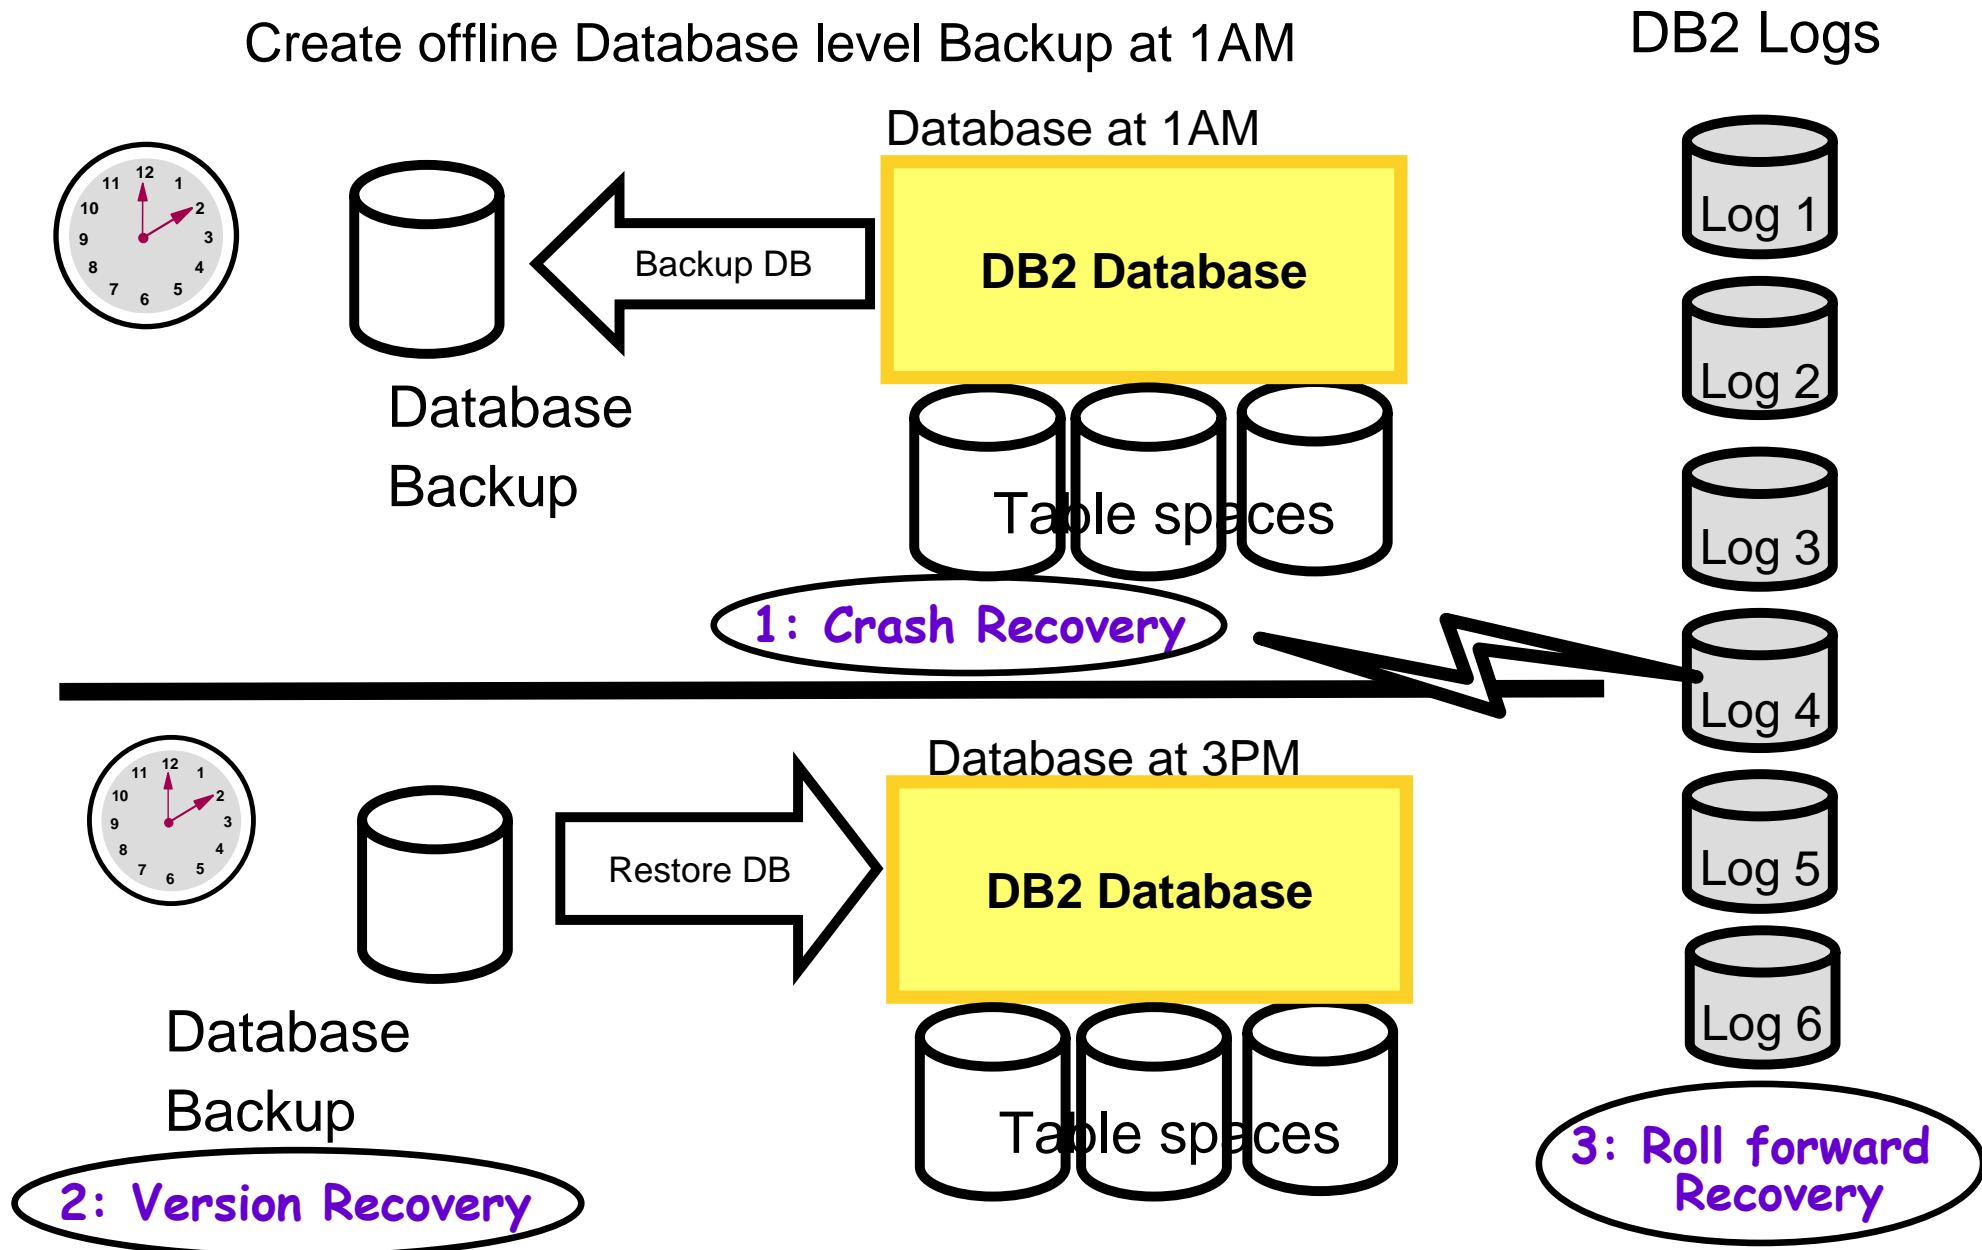


Unit objectives


After completing this unit, you should be able to:

- Describe the major principles and methods for backup and recovery
- State the three types of recovery used by DB2
- Explain the importance of logging for backup and recovery
- Describe how data logging takes place, including circular logging and archival logging
- Use the BACKUP, RESTORE, ROLLFORWARD and RECOVER commands
- Perform a table space backup and recovery
- Restore a database to the end of logs or to a point-in-time
- Discuss the configuration parameters and the recovery history file and use these to handle various backup and recovery scenarios

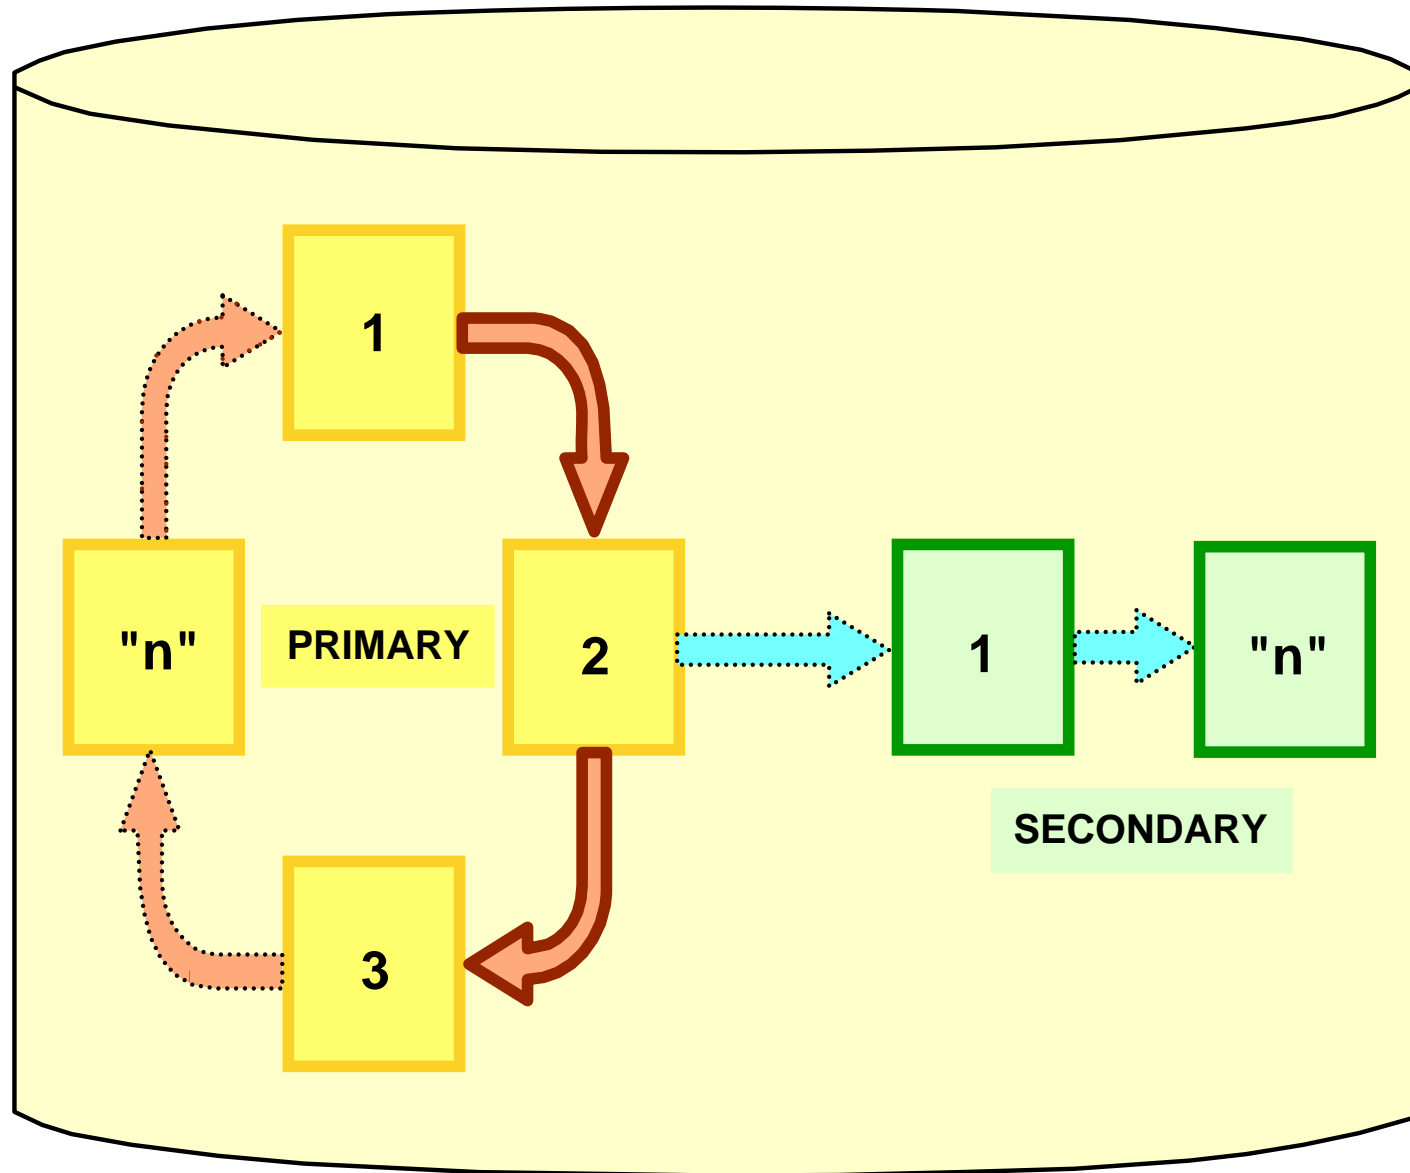
DB2 Database recovery methods



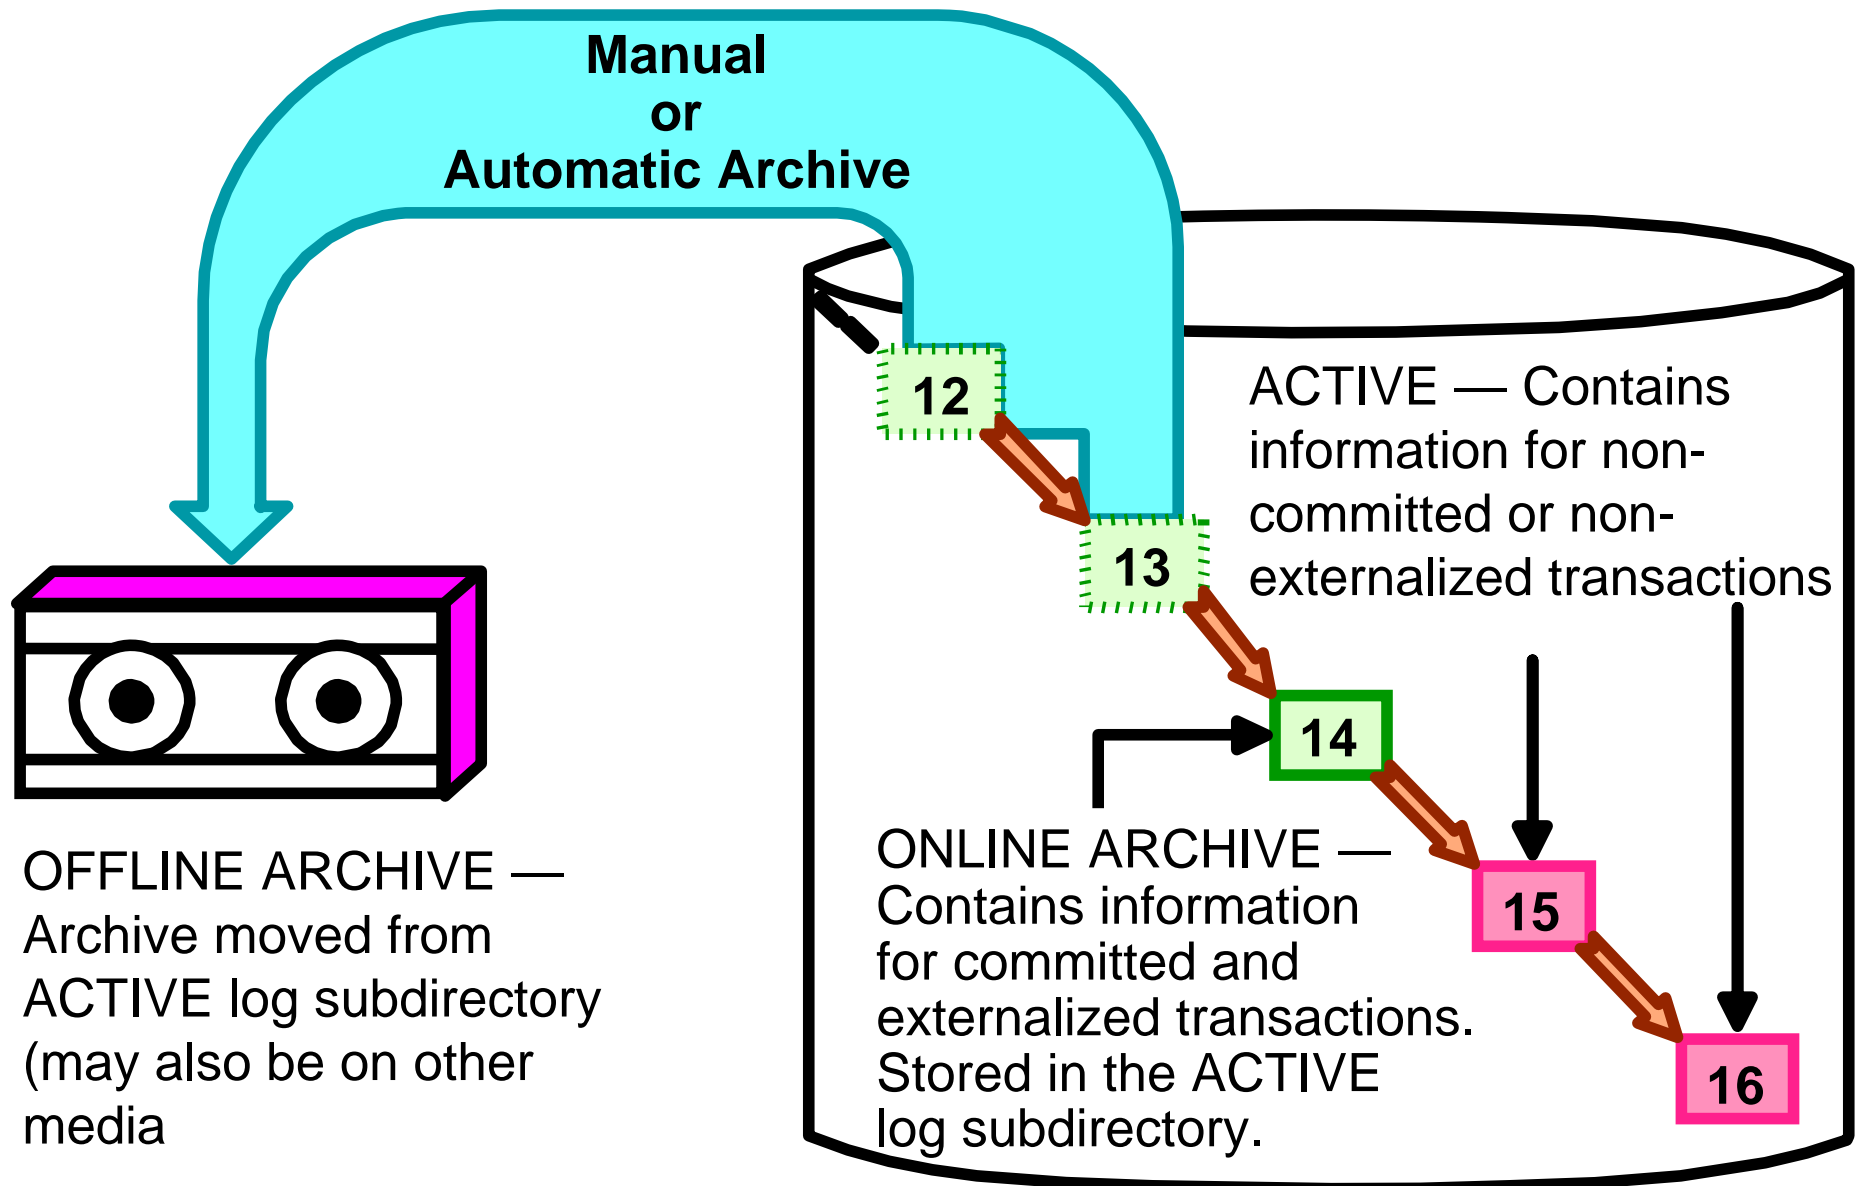
Introduction to logging



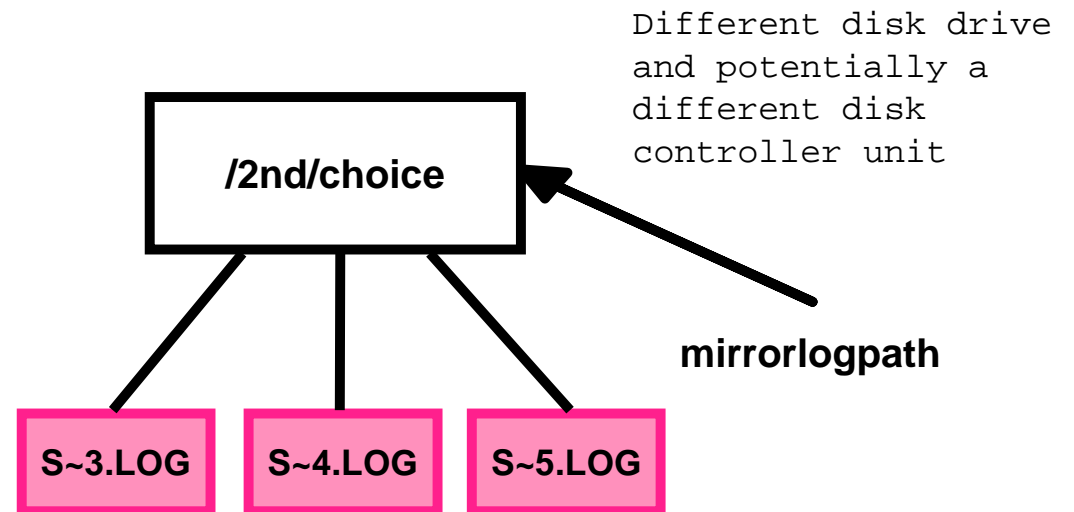

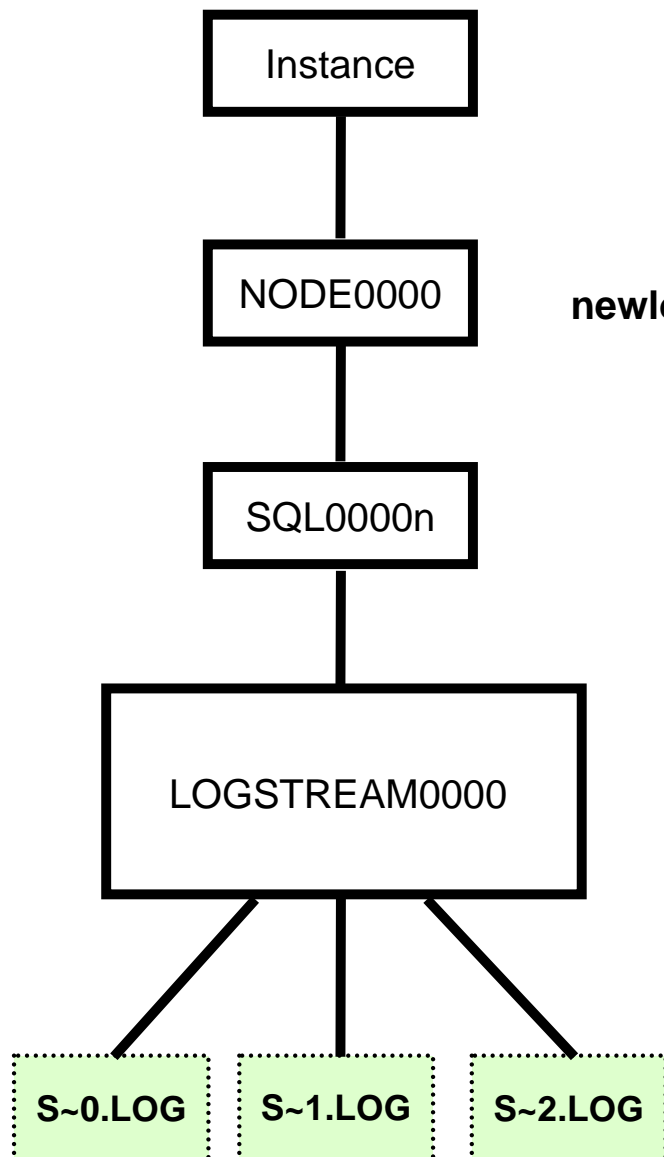
Circular logging (Non-recoverable database)



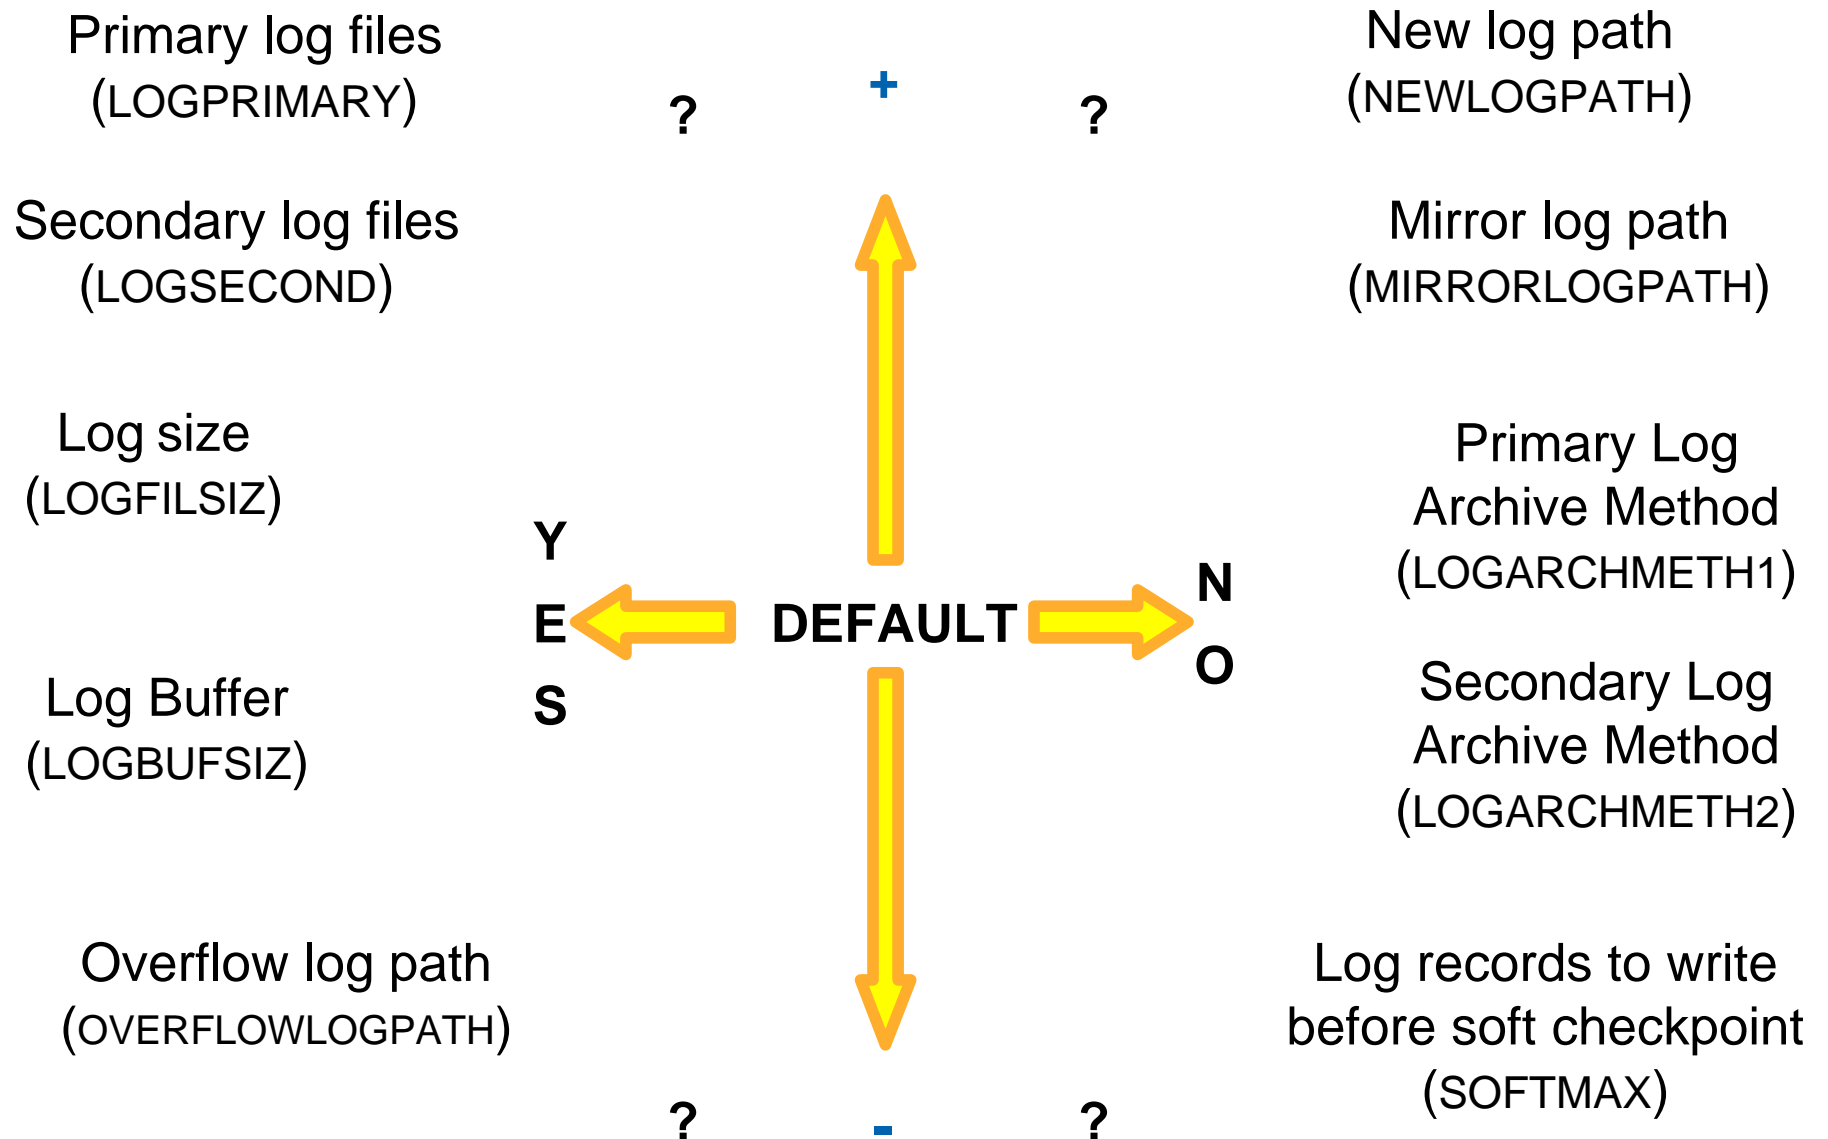
Archival logging (Recoverable database)



Location of log files



Configure database logs: Parameters



DB2 recovery-related system files

- Recovery History File - db2rhist.asc:
 - Created during Create Database command
 - Updated by DB2 Utility execution:
 - Back up database or table space
 - Restore database or table space
 - Roll forward database or table space
 - Load table
 - Reorg table
 - DB2 Log file archival
 - Create/Alter/Rename table space
 - Quiesce table spaces for table
 - Drop Table (optional)
 - Included on each DB2 backup file
 - db2 LIST HISTORY command
- Log Control File – SQLOGCTL.LFH.n and SQLOGCTL.GLFH.n:
 - Used during crash recovery
 - Disk updated at end of each Log File
 - Use softmax value < 100 (% of logfilsiz) to refresh pointers more often
 - Included at end of each DB2 backup

Backup Utility options (partial)

BACKUP DATABASE database-alias [USER username [USING password]]

[TABLESPACE (tblspace-name [{,tblspace-name} ...])] [ONLINE]

[INCREMENTAL [DELTA]]

[USE {TSM | XBSA} [OPEN num-sess SESSIONS]

[OPTIONS {options-string | options-filename}] | TO dir/dev

[{,dir/dev} ...] | LOAD lib-name [OPEN num-sess SESSIONS]
[OPTIONS {options-string | options-filename}]]

[WITH num-buff BUFFERS] [BUFFER buffer-size] [PARALLELISM n]

[COMPRESS [COMPRlib lib-name [EXCLUDE]] [COMPROPTS options-string]]

[UTIL_IMPACT_PRIORITY [priority]
[{INCLUDE | EXCLUDE} LOGS] [WITHOUT PROMPTING]


The backup files

- File name for backup images *on disk* has:
 - Database alias
 - Type of backup (0=full, 3=table space, 4=copy from table load)
 - Instance name
 - Database Partition (DBPARTnnn, nnn is 000 for single partition DB)
 - Timestamp of backup (YYYYMMDDHHMMS)
 - Sequence number (for multiple files)




MUSICDB.0.DB2.DBPART000. 20120522120112 .001

- Tape images are not named, but internally contain the same information in the backup header for verification purposes
- Backup history provides key information in easy-to-use format

Syntax of the RESTORE command (partial)



Restore options:



Backup/restore table space considerations

- Recoverable Database - Roll-forward must be enabled
- Can choose to restore a subset of table spaces
- Generally best to put multiple spaces in one backup image:
 - Makes table space recovery strategy easier
 - Provides access to related tables spaces and coherent management of these table spaces
- Handling of long/LOB/XML data requires a correlated strategy
- Point-in-time recovery is supported, but has requirements
- Faster recovery for Catalogs using Tablespace Level backup
- Critical business application tables should obviously be the *focus* of the backup/restore, but other tables are needed in support of these tables

Roll forward pending state

- *Roll forward pending* is set as a result of:
 - Restore of *offline* database backup omitting the command option WITHOUT ROLLING FORWARD
 - Restore of an *online* database backup
 - Restore of *any table space* level backup
 - DB2 detects media failure isolated at a table space
- Scope of pending state managed by DB2:
 - *Database* in pending state will not permit any activity
 - *Table spaces* in pending state will permit access to other table spaces

Syntax of the ROLLFORWARD command

```
>>-ROLLFORWARD--+--DATABASE--+--database-alias----->
                '-DB-----'
```

```

>--+-----+
+-->
|               .-ON ALL DBPARTITIONNUMS-.   .-USING UTC TIME---.
+--TO--+--isotime--+-----+-----+-----+-----+-----+-----+
|               |               '-USING LOCAL TIME-' | +-AND COMPLETE--+
|               |               .-ON ALL DBPARTITIONNUMS-. | '-AND STOP-----'
|               +-END OF BACKUP+-----+-----+-----+
|               '-END OF LOGS--+-----+-----+-----+'
|               '-| On Database Partition clause |-'
|
+-+--COMPLETE-----+-----+-----+-----+-----+-----+
+-STOP-----+-----+-----+-----+-----+     '-| On Database Partition clause |-'
+-CANCEL-----+-----+-----+-----+-----+
|               .-USING UTC TIME---. |
|               '-QUERY STATUS--+-----+-----+-----+'
|               '-USING LOCAL TIME-'

```

```
>-----+----->  
'-TABLESPACE-+-ONLINE-  
|             |  
|      .-'---'-.  
|      v         |  
|'-((---tablespace-name-+-))--+-+-----+-'  
|                                     '-ONLINE-'
```

[illegible]

```
>----->
'-NORETRIEVE-'
```


ROLLFORWARD: How far?

- END OF LOGS: (Apply as many changes as possible):
 - Rollforward will apply all available logs beginning with the logs associated with the backup that was restored
 - Archived logs will be retrieved unless NORETRIEVE is specified
- Point-in-time (PIT): (Apply changes up to a specified time):
 - Specified in Coordinated Universal Time (UTC) via command
 - Specified in local time *on server* with USING LOCAL TIME
 - Specified in local time on the client via GUI interface
 - Format: yyyy-mm-dd-hh.mm.ss.nnnnnn
- END OF BACKUP: (Apply as few changes as possible):
 - Allows a Database to be recovered from an online database backup and to end the ROLLFORWARD processing at the earliest point where the database is consistent.
 - Recovery history file (RHF) shows logs associated with online backups
- Table space point-in-time considerations:
 - Minimum roll forward time maintained for each table space – requires roll forward at least to the last DDL change (create, alter, drop) in a table space
 - Table spaces are placed in backup pending when the roll forward completes to insure future recoverability

DB2 RECOVER command

```
db2 recover database salesdb
```

- RECOVER command performs both RESTORE and ROLLFORWARD command processing using Recovery History file data.
- Can use Full or Incremental Database level backup images
- Allows Partitioned DPF database to be recovered using a single command
- Database can be recovered to *End of logs* or a point in time.
- RESTART option forces failed RECOVER command to redo RESTORE even if previous restore completed
- USING HISTORY FILE option allows disaster recovery with no existing database
- No table space level recovery options

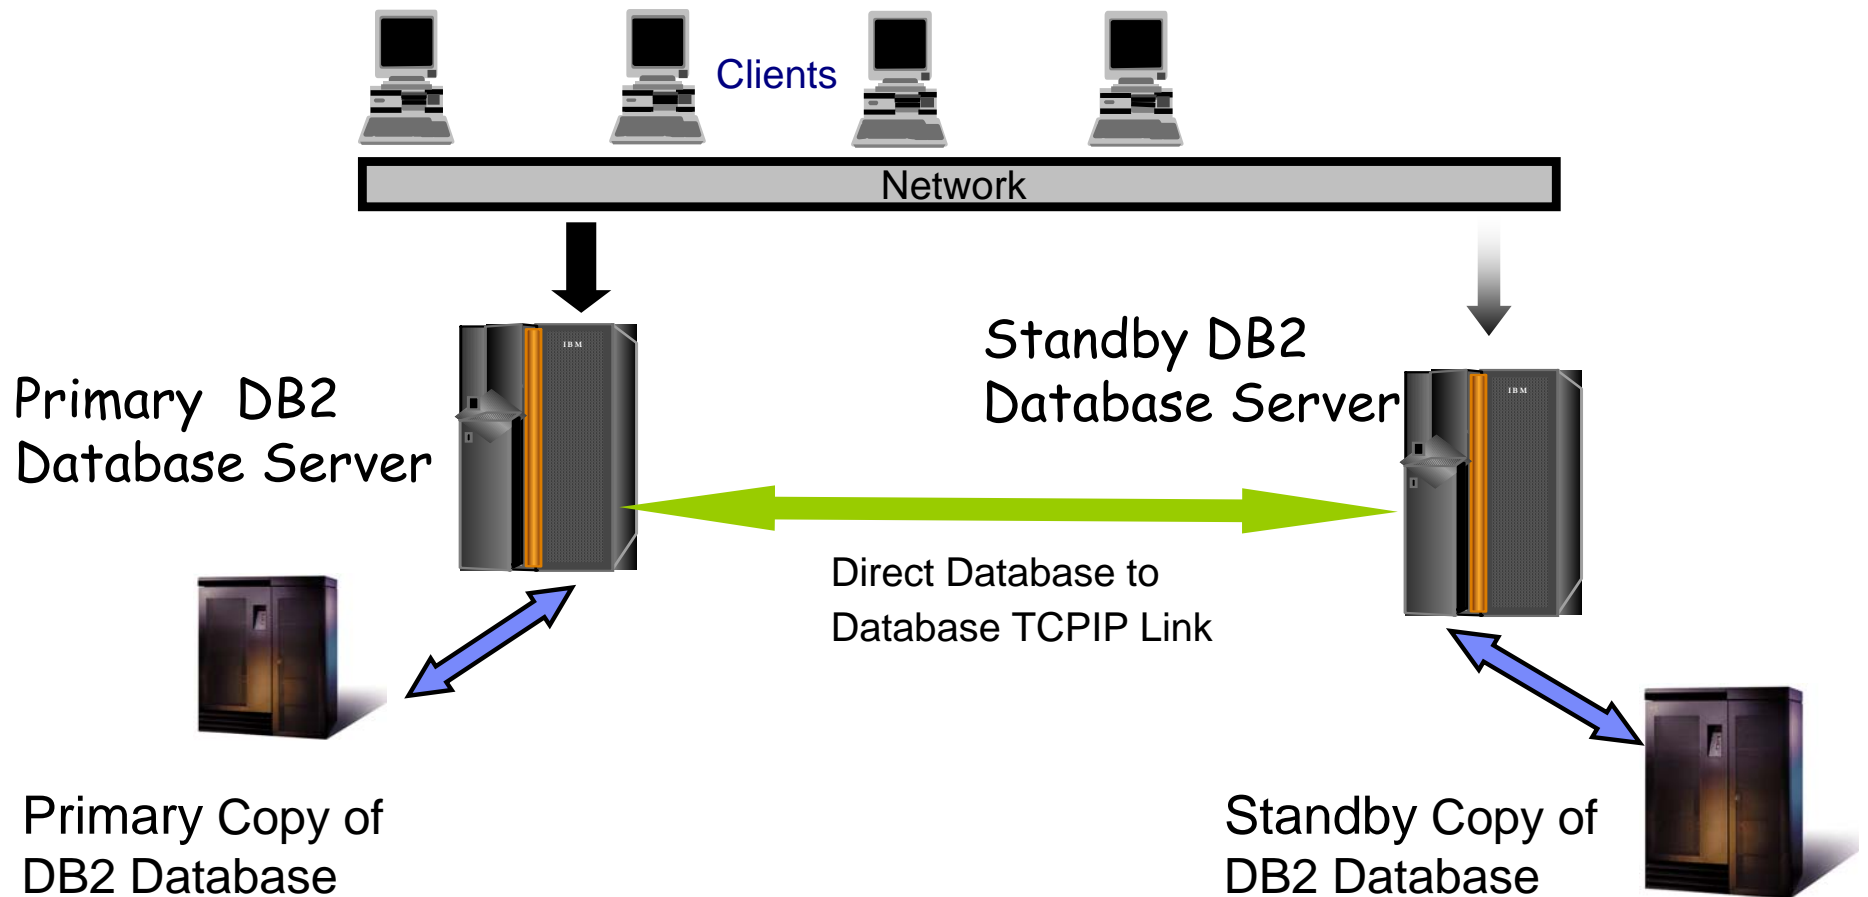


Disaster recovery considerations

Should you use database recovery, table space recovery, or a combination?


- Database level backups:
 - Offline backup does not require a roll forward
 - Online backup will include the 'required' logs for a roll forward TO END OF BACKUP
 - Entire database can be restored on another system
 - Reduces backup management
- Table space level backups:
 - Can be used to supplement database backups
 - A complete set of tablespace backups could be used to REBUILD a database for disaster recovery
 - Increases the number of backup images
- Using the REBUILD WITH option of RESTORE simplifies creating a full or partial copy of a database using either database or table space backup images

High Availability Disaster Recovery: HADR



High Availability Disaster Recovery (HADR) feature is included in all the for-purchase DB2 10.1 editions (Express, Workgroup Server, Enterprise Server, and Advanced Enterprise Server, Database Enterprise Developer).

DB2 10.1 support for multiple active standby databases



- HADR feature in multiple standby mode allows up to three standby databases to be configured
- One Standby is designated the *principal HADR standby database*
- Any additional standby database is an *auxiliary HADR standby database*
- Both types of HADR standbys:
 - Are synchronized with the HADR primary database through a direct TCP/IP connection
 - Support reads on standby
 - Can issue a forced or non-forced takeover
- Other HADR enhancements included in DB2 10.1
 - Log spooling on the Standby database
 - Delayed replay for a Standby database

DB2 Integrated Cluster Manager support

- HA Cluster Manager Integration:
 - Coupling of DB2 and TSA (S AMP) on Linux, Solaris and AIX.
 - TSA can be installed and maintained with DB2 installation procedures.
 - DB2 interface to configure cluster.
 - DB2 to maintain cluster configuration, add node, add table space, and so on. Exploitation of new vendor independent layering (VIL), providing support for any cluster manager.
- NO SCRIPTING REQUIRED!
 - One set of embedded scripts that are used by all cluster managers.
- Automates HADR failover
 - Exploit HA cluster manager integration to automate HADR Takeover on Standby.

DB2: Cluster configuration simplified

- DB2 utility db2haicu:
 - DB2 High Availability Instance Configuration Utility
 - Sets up DB2 with the Cluster Manager
 - For HADR or Shared Storage cluster
 - Interactive or XML file driven interface
 - Sets the DBM CFG option **cluster_mgr**
- DB2 communicates with the Cluster Manager:
 - Keeps the Cluster Manager in sync with DB2 changes:
 - CM needs to be aware of new or changed table spaces, containers, and so on
 - Avoids failover problems due to missing resources
 - Keeps the XML cluster configuration file up to date
 - Automates HADR failover
- Check IBM Developerworks for articles that provide detailed examples and suggestions for DB2 Cluster configuration

DB2 additional recovery facilities

- On-demand log archiving
- Infinite active logs
- Block transactions on log directory full
- Split mirror database copies:
 - SET WRITE SUSPEND/RESUME commands
 - db2inidb command modes:
 - **SNAPSHOT**: Database copy for testing reporting
 - **STANDBY**: Database copy to create standby database for quick recovery
 - **MIRROR**: Use split mirror database copy instead of RESTORE
- Incremental and delta database and table space backups
- Relocating a database or a table space:
 - RESTORE UTILITY with REDIRECT option
 - db2relocatedb command
- Full and partial database REBUILD support
- Integrated Cluster Failover support using TSA

DB2 for LUW Advanced Database Recovery

CL492: *DB2 for LUW Advanced Database Recovery (4 days with Labs)*:

- Automated Recovery log management
- Dropped table recovery
- Table space recovery
- Incremental backup and restore
- Database crash recovery
- Database and table space relocation methods
- Additional recovery facilities
- Recovery history file
- Disaster recovery of DB2 databases
- DB2 high availability and split mirror support
- DB2 partitioned database recovery considerations
- DB2 high availability disaster recovery (HADR) implementation

Unit summary

Having completed this unit, you should be able to:

- Describe the major principles and methods for backup and recovery
- State the three types of recovery used by DB2
- Explain the importance of logging for backup and recovery
- Describe how data logging takes place, including circular logging and archival logging
- Use the BACKUP, RESTORE, ROLLFORWARD and RECOVER commands
- Perform a table space backup and recovery
- Restore a database to the end of logs or to a point-in-time
- Discuss the configuration parameters and the recovery history file and use these to handle various backup and recovery scenarios

Student exercise

